

class07QD

Ilyas Darif A16577084

2024-04-23


today we will start our multi part exploration of some key machine learning methods. we will begin with clustering - finding groupings in data and then dimensionality reduction

Clustering

lets start with “k-means” clustering the main function in base R for the is `means()`

```
#makeup up some data  
hist( rnorm(100000, mean=3))
```

Histogram of `rnorm(1e+05, mean = 3)`



```
tmp <- c(rnorm(30, -3), rnorm(30, +3))
```


```
x <- cbind(x=tmp, y=rev(tmp))
```

```
x
```

	x	y
[1,]	-2.6113127	5.1299379
[2,]	-4.1261385	4.4627565
[3,]	-2.5490062	4.0816916
[4,]	-3.8114809	1.9800675
[5,]	-2.8344824	3.0857448
[6,]	-4.0974978	2.9505656
[7,]	-3.3632147	1.3392605
[8,]	-2.3642414	4.1934734
[9,]	-2.8421212	2.8643826
[10,]	-2.5066374	3.2811309
[11,]	-2.5657316	2.6727083
[12,]	-2.2222934	3.5793545
[13,]	-1.5879509	2.3166419
[14,]	-2.7111841	3.7354913
[15,]	-2.2686093	1.7676810
[16,]	-3.9412907	2.6022936
[17,]	-2.4557176	2.9250352
[18,]	-4.1346654	2.8126393
[19,]	-4.0804667	3.6907071
[20,]	-3.5474158	0.9278268
[21,]	-1.8138030	2.2516766
[22,]	-3.1880890	4.0844581
[23,]	-2.5697118	2.6233017
[24,]	-3.2722725	2.4793893
[25,]	-3.1490273	4.3870648
[26,]	-2.9280196	4.3700703
[27,]	-3.5594731	3.9197163
[28,]	-2.9141401	4.0696375
[29,]	-2.5886518	3.1934598
[30,]	-3.6692415	3.8313418
[31,]	3.8313418	-3.6692415
[32,]	3.1934598	-2.5886518
[33,]	4.0696375	-2.9141401
[34,]	3.9197163	-3.5594731
[35,]	4.3700703	-2.9280196
[36,]	4.3870648	-3.1490273

```
[37,] 2.4793893 -3.2722725
[38,] 2.6233017 -2.5697118
[39,] 4.0844581 -3.1880890
[40,] 2.2516766 -1.8138030
[41,] 0.9278268 -3.5474158
[42,] 3.6907071 -4.0804667
[43,] 2.8126393 -4.1346654
[44,] 2.9250352 -2.4557176
[45,] 2.6022936 -3.9412907
[46,] 1.7676810 -2.2686093
[47,] 3.7354913 -2.7111841
[48,] 2.3166419 -1.5879509
[49,] 3.5793545 -2.2222934
[50,] 2.6727083 -2.5657316
[51,] 3.2811309 -2.5066374
[52,] 2.8643826 -2.8421212
[53,] 4.1934734 -2.3642414
[54,] 1.3392605 -3.3632147
[55,] 2.9505656 -4.0974978
[56,] 3.0857448 -2.8344824
[57,] 1.9800675 -3.8114809
[58,] 4.0816916 -2.5490062
[59,] 4.4627565 -4.1261385
[60,] 5.1299379 -2.6113127
```

```
plot(x)
```



now lets try out kmeans()

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	3.186984	-3.009130
2	-3.009130	3.186984

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 43.05918 43.05918
(between_SS / total_SS = 93.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
[1] "kmeans"
```

Q1. how many points in each cluster?

```
km$size
```

```
[1] 30 30
```

Q2. what component of your result object details cluster assignment/membership

```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Q3. what are centers/mean values of each cluster

```
km$centers
```


```
      x      y
1 3.186984 -3.009130
2 -3.009130  3.186984
```

Q4. make a plot of your data showing your clustering results (groupings/clusters and cluster centers)

```
plot(x, col=c("red", "blue"))
```




```
plot(x, col=c(1,2))
```




```
plot(x, col=km$cluster)

points(km$centers, col="green", pch=15, cex=3)
```



Q5. run `kmeans()` again and cluster in 4 groups and plot the results.

```
km4 <- kmeans(x, centers = 4)
plot(x, col=km4$cluster)
```



hierarchial clustering

this form of clustering aims to reveal the structure in your data by progressively grouping points into a ever smaller number of clusters

the maion function in base R for this called `hclust()` . this function does not take our input data directly but wants a “distance matrix” that details how (dis)similar our inout points are to each other


```
hc <- hclust( dist(x) )
hc
```

Call:

```
hclust(d = dist(x))
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

the print out above is not useful (unlike that from `kmeans`) but there is a useful `plot()` method



Principal Component Analysis (PSA)

the goal of PCA is to reduce the dimensionality of a dataset down to some smaller subset of new variables (called PCs) that are a useful bases for further analysis, like visualization, clustering ect

Q1.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143

Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41


```
dim(x)
```

```
[1] 17  4
```

Q2. the `row.names = 1` way because it was a little more simple to use


Q3.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```




to make the plot the other bar style you change `beside=T` to `beside=F`

```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q4(5). it means these are the axis for each plot

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



so the paris plot is useful for small datasets but it can be lots of work to interpret and gets interactable for longer datasets

So PCA to the rescue... the main function to do PCA in base R is called `prcomp()`. this function wants the transpose of our data in this case.

`t(x)`

	Cheese	Carcass_meat	Other_meat	Fish	Fats_and_oils	Sugars
England	105	245	685	147	193	156
Wales	103	227	803	160	235	175
Scotland	103	242	750	122	184	147
N.Ireland	66	267	586	93	209	139
	Fresh_potatoes	Fresh_Veg	Other_Veg	Processed_potatoes		
England		720	253	488		198
Wales		874	265	570		203
Scotland		566	171	418		220
N.Ireland		1033	143	355		187
	Processed_Veg	Fresh_fruit	Cereals	Beverages	Soft_drinks	
England		360	1102	1472	57	1374
Wales		365	1137	1582	73	1256
Scotland		337	957	1462	53	1572

N.Ireland	334	674	1494	47	1506
	Alcoholic_drinks	Confectionery			
England	375		54		
Wales	475		64		
Scotland	458		62		
N.Ireland	135		41		

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

```
[1] "prcomp"
```


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

A MAJOR pcA result viz is called a “PCA plot” (aka a score plot, biplot, pc1 vs pc2 plot, ordination plot)

```
mycols <- c("orange", "red", "blue", "darkgreen")
plot(pca$x[,1], pca$x[,2], col=mycols, pch=16, xlab="PC1", ylab="PC2")
abline(h=0, col="gray")
```

```
abline(v=0, col="gray")
```



another important output from PCA is called the “loadings” vector or the “rotation” component - this tells us how much the original variables (the foods in this case)

```
pca$rotation
```

	PC1	PC2	PC3	PC4
Cheese	-0.056955380	0.016012850	0.02394295	-0.694538519
Carcass_meat	0.047927628	0.013915823	0.06367111	0.489884628
Other_meat	-0.258916658	-0.015331138	-0.55384854	0.279023718
Fish	-0.084414983	-0.050754947	0.03906481	-0.008483145
Fats_and_oils	-0.005193623	-0.095388656	-0.12522257	0.076097502
Sugars	-0.037620983	-0.043021699	-0.03605745	0.034101334
Fresh_potatoes	0.401402060	-0.715017078	-0.20668248	-0.090972715
Fresh_Veg	-0.151849942	-0.144900268	0.21382237	-0.039901917
Other_Veg	-0.243593729	-0.225450923	-0.05332841	0.016719075
Processed_potatoes	-0.026886233	0.042850761	-0.07364902	0.030125166
Processed_Veg	-0.036488269	-0.045451802	0.05289191	-0.013969507
Fresh_fruit	-0.632640898	-0.177740743	0.40012865	0.184072217
Cereals	-0.047702858	-0.212599678	-0.35884921	0.191926714

Beverages	-0.026187756	-0.030560542	-0.04135860	0.004831876
Soft_drinks	0.232244140	0.555124311	-0.16942648	0.103508492
Alcoholic_drinks	-0.463968168	0.113536523	-0.49858320	-0.316290619
Confectionery	-0.029650201	0.005949921	-0.05232164	0.001847469

PCA looks to be a super useful method for gaining some insight into high dimensional data that is difficult to examine in other ways